

CPU Time = CPU Clock Cycles \* Clock Cycle Time = CPU Clock Cycles / Clock Rate  
 Clock rate = Clock cycles / CPU time  
 Clock cycles = instruction count \* cycle per instruction = (instruction count \* CPI) / Clock Rate  
 Clock Cycles =  $\sum_{i=1}^n CPI_i \cdot Instruction\ Count_i$   
 Weighted average CPI =  $\sum_{i=1}^n CPI_i \cdot (Instruction\ Count_i / instruction\ Count)$   
 CPU Time = (Instructions / program) \* (Clock Cycles / instructions) \* (seconds / clock cycle)  
 T improved = T affected / improvement factor + T unaffected

```

A[5] = h + A[3]
while (save[i] == k) i += 1;
loop: slt $t1, $s3, 2
      add $t1, $t1, $s6
      lw $t0, 0($t1)
      bne $t0, $s3, exit
      addi $s3, $s3, 1
      sw $t0, 20($s3)
      exit: j loop
  
```

OP	RS	RT	RD	SHAMT	FUNCT
6	5	5	5	5	6

Jump also in I-format

OP	RS	RT	CONSTANT/ADDRESS
6	5	5	16

0	0000	4	0100	8	1000	c	1100
1	0001	5	0101	9	1001	d	1101
2	0010	6	0110	a	1010	e	1110
3	0011	7	0111	b	1011	f	1111

正的到负的要记得专为相反的 再加一

fact:  
 addi \$sp, \$sp, -8 # adjust stack for 2 items  
 sw \$ra, 4(\$sp) # save return address  
 sw \$a0, 0(\$sp) # save argument  
 slti \$t0, \$a0, 1 # test for n < 1  
 beq \$t0, \$zero, L1  
 C code:  
 int fact (int n)  
 {  
 if (n < 1) return 1;  
 else return n \* fact(n - 1);  
 }  
 Argument n in \$a0  
 Result in \$v0

addi \$v0, \$zero, 1 # if so, result is 1  
 addi \$sp, \$sp, 8 # pop 2 items from stack  
 jr \$ra. # and return

L1:  
 addi \$a0, \$a0, -1 # else decrement n  
 jal fact # recursive call  
 lw \$a0, 0(\$sp) # restore original n  
 lw \$ra, 4(\$sp) # and return address  
 addi \$sp, \$sp, 8 # pop 2 items from stack  
 mul \$v0, \$a0, \$v0 # multiply to get result  
 jr \$ra # and return

lb t0ff(b) # \$t ← Sign-extended byte  
 lbu t0ff(b) # \$t ← Zero-extended byte  
 sb t0ff(b) # The byte at offset ← low-order # byte from register \$t.

lh t0ff(b) # \$t ← Sign-extended halfword  
 lhu t0ff(b) # \$t ← zero-extended halfword  
 sh t0ff(b) # Halfword at offset ← low-order # two bytes from \$t.

lui Copies 16-bit constant to left 16 bits of rt and Clears right 16 bits of rt to 0

Dynamic linking ... Requires procedure code to be relocatable & Automatically picks up new library versions


```

swap: slt $t1, $a1, 2 # $t1 = k * 4
      add $t1, $a0, $t1 # $t1 = v + (k * 4) # _address of v[k]
      lw $t0, 0($t1) # $t0 (temp) = v[k]
      lw $t2, 4($t1) # $t2 = v[k+1]
      sw $t2, 0($t1) # v[k] = $t2 (v[k+1])
      sw $t0, 4($t1) # v[k+1] = $t0 (temp)
      jr $ra # return to calling routine
  
```

### chapter 3 a)


- NOR: is a 0 if either operand is a 1
- NAND: is a 0 only if both operands are 1
- XOR: is a 1 if the operands are different

pos



To compute slt A B:

subtract B from A (set binvert and the LS Carry In to 1)  
 Result for all 1-bit ALUs except the LS should always be 0  
 Result for the LS 1-bit ALU should be the result bit from the MS 1-bit ALU  
 LS: Least significant (rightmost) MS: Most significant (leftmost)



```

avg:  move $t0, $a0
      move $t1, $a1
      li $t2, 0
      la $t3, const0
      lwc1 $f16, 0($t3)
      beq $t2, $t1, done
      add.s $f17, 0($t0)
      addi $t2, $t2, 1
      addi $t0, $t0, 4
      j loop
      done: la $t3, const12
           lwc1 $f17, 0($t3)
           div.s $f0, $f16, $f17
           jr $ra
  
```

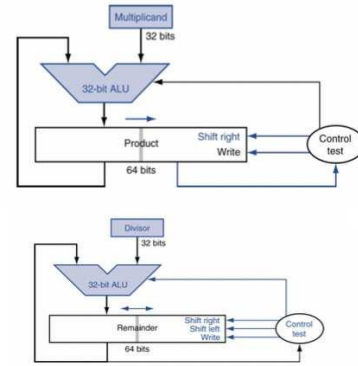
c.lt.s \$f12, \$f13  
 beqz w.anchorage  
 la \$a0, troy\_str  
 li \$v0, 4  
 syscall  
 j w\_done

$$C_{i+1} = G_i + P_i \cdot C_i$$


$$G_i = A_i \cdot B_i$$

$$P_i = A_i \oplus B_i$$

Operation	Operand A	Operand B	overflow
X = A + B	A ≥ 0	B ≥ 0	X < 0
X = A + B	A < 0	B < 0	X ≥ 0
X = A - B	A ≥ 0	B < 0	X < 0
X = A - B	A < 0	B ≥ 0	X ≥ 0



div \$a0, \$a1  
 • mfhi \$a2 # remainder to \$a2  
 • mflo \$v0 # quotient to \$v0



single: 8 bits  
 double: 11 bits

S	Exponent	Fraction
---	----------	----------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0 ⇒ non-negative, 1 ⇒ negative)
- Normalized significand:  $1.0 \leq |\text{significand}| < 2.0$ 
  - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
  - Significand is Fraction with the "1." restored
- Exponent: actual exponent + Bias
  - Ensures exponent is unsigned
  - Single: Bias = 127; Double: Bias = 1023

Single precision 32 bit ... 6 decimal digits

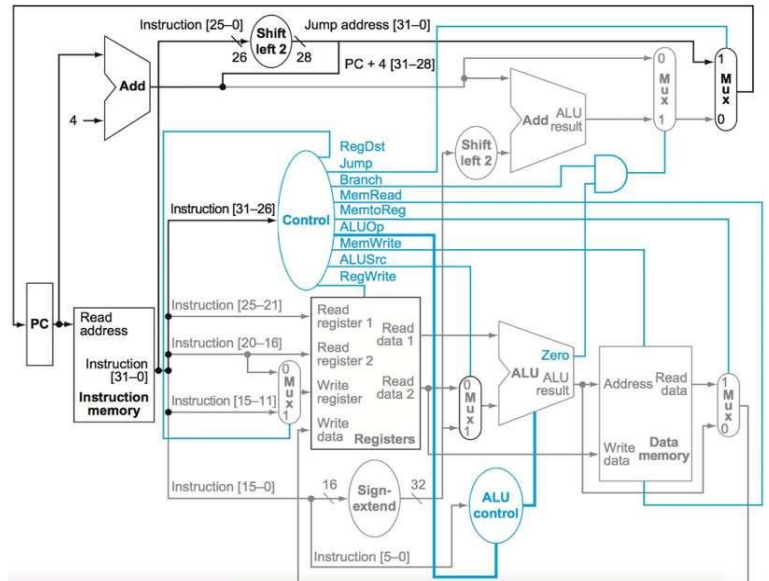
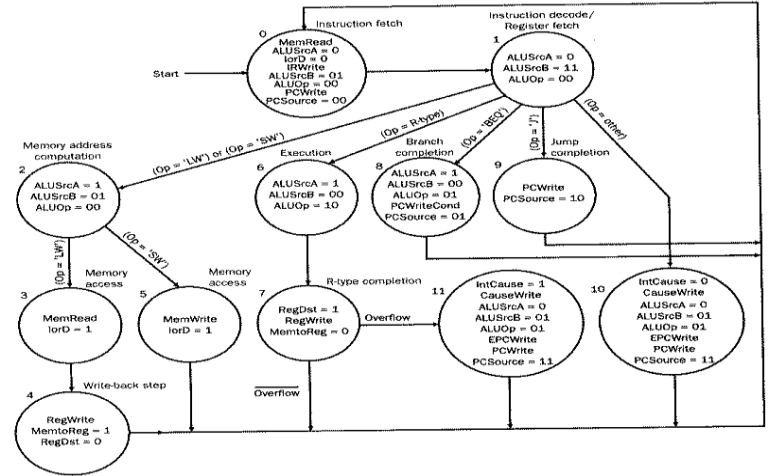
Double precision 64 bit ... 16 decimal digits

Binary ... dividing right ... multiplying left

Single- and double-precision comparison

e.g. c.xx.s, c.xx.d (xx is eq, lt, le, ...) with bclt & bclt

C code:  
 float f2c (float fahr) {  
 return ((5.0/9.0)\*(fahr - 32.0));  
 }  
 • fahr in \$f12, result in \$f0, literals in global memory space  
 Compiled MIPS code:  
 f2c: lwc1 \$f16, const5(\$gp)  
 lwc1 \$f18, const9(\$gp)  
 div.s \$f16, \$f16, \$f18  
 lwc1 \$f18, const32(\$gp)  
 sub.s \$f18, \$f12, \$f18  
 mul.s \$f0, \$f16, \$f18  
 jr \$ra



chapter 4 b)

```
lw $1, addr, ...MEM
add $4, $5, $6 ... EX
beq stalled
beq $1, $4, target ...pointing to ID

lw $1, addr, ... MEM
beq stalled
beq $1, $0, target ... pointing to ID
```

2: Branch Target Buffer (BTB)

Cache that stores: the PCs of branches  
the predicted target address  
branch prediction bits

Accessed by PC address in fetch stage  
if hit: address was for this branch instruction  
fetch the target instruction if prediction bits say taken

No branch delay if: branch found in BTB  
prediction is correct  
(assume BTB update is done in the next cycles)

A single prediction bit does not work well with loops

- mispredicts the first & last iterations of a nested loop

Two-bit branch prediction for loops

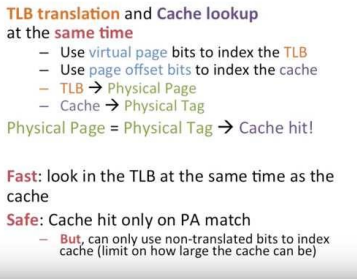
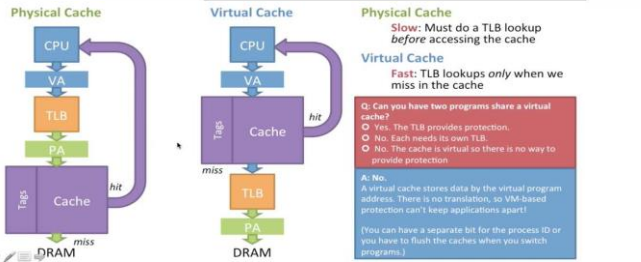
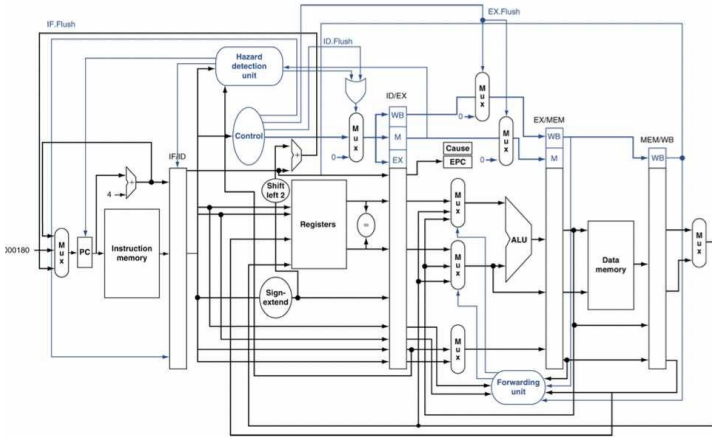
- Algorithm: have to be wrong twice before the prediction is changed

Exception ... CPU ... division by zero and undefined ... system control coprocessor ... save PC  
offending ... mips for Exception Program Counter  
Eg. add \$1, \$2, \$1

- o Flush add and subsequent instructions
- o Set Cause and EPC register values
- o Transfer control to handler

Interrupt ... I/O controller ... keyboard

Multiple exception ... overlap ... flush subsequent instruction & precise exception  
Imprecise exceptions ... stop & handler  
Instruction-level parallelism ... deeper pipeline (shorter clock cycle) & multiple issue (static ...  
made by the compiler before execution ... issue packets & dynamic ... during execution by the CPU)



Quiz summary

当新建 array 的时候在 heap 上, function 传入的时候要加& (地址传入)  
Echo \$? Last process call  
0 successfully 1 unsuccessful

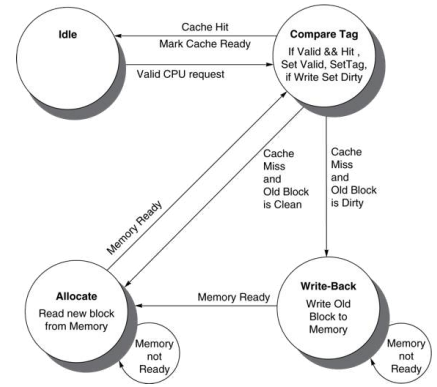
```
addi $sp, $sp, -12
sw $s0, 8($sp)
sw $s1, 4($sp)
sw $s2, 0($sp)

lw $ra, 0($sp)
addi $sp, $sp, 4
```

For memory operations on full words, an address aligned to the word boundary (i.e., divisible by 0x04) is required. The instruction in question attempts to access memory specifying the effective address which ends in 0x01, i.e., which is not aligned.

Least significant right-most  
Most significant left-most

Functionally complete set  
Except XOR AND OR



Chapter 5 a) & b)

Recently disk to DRAM  
MORE recently DRAM to SRAM

Hit:  
Write-Through  
Cache immediately writes modified block to memory ... write buffer  
Write-back (VM)  
先写到 cache, writes back to memory when the block is evicted ... dirty bit  
Miss:  
Write allocation  
Read block into the cache, update the copy in the cache

stall cycles = (num of reads / num of program) \* read miss rate \* read miss penalty  
stall cycles = (num of writes / num of program) \* write miss rate \* write miss penalty + write buffer stalls

AMAT = Hit time + Miss rate \* Miss penalty  
Miss penalty = access time / (1/clock rate)

Page fault ... LRU ... reference bit set to 1 -> periodically cleared to 0 -> reference bit to 1 (not used recently)

Compulsory misses (aka cold start misses)

First access to a block  
Capacity misses  
Due to finite cache size  
A replaced block is later accessed again  
Conflict misses (aka collision misses)  
In a non-fully associative cache  
Due to competition for entries in a set  
Would not occur in a fully associative cache of the same total size

Cache Coherence Protocols

Migration of data to local caches  
o Reduces bandwidth for shared memory  
Replication of read-shared data  
o Reduces contention for access  
Snooping  
o Writes broadcast on shared bus  
Directory-based  
o Each block assigned an ordering point

Use C system calls:

- void \*malloc(size\_t size) returns a pointer to a chunk of memory which is the size in bytes requested  
- void \*calloc(size\_t nmem, size\_t size) same as malloc but puts zeros in all bytes and asks for the number of elements and size of an element.  
- void free(void \*ptr) deallocates a chunk of memory. Acts much like the delete operator.  
- void \*realloc(void \*ptr, size\_t size) changes the size of the memory chunk pointed to by ptr to size bytes.

~ current directory

Absolute path starting with /  
-ls lists file names (like DOS dir command).  
-who lists users currently logged in.  
-date shows the current time and date.  
-pwd print working directory

. for current, . for parent

df: shows what disk holds a directory.

cp [options] source dest //copy

rm [options] names... //remove

- when the file was last changed: ls -l

- when the file was created\*: ls -lc

- when the file was last accessed: ls -ul

Files:

-r: allowed to read.

-w: allowed to write.

-x: allowed to execute

• Directories:

-r: allowed to see the names of the files.

-w: allowed to add and remove files.

-x: allowed to enter the directory

chmod [ugoalrwx] file

ls \* //all files and sub files ... ls -al \$HOME

ls -al \* // including the directory

ls a\* // starting with a

ls \*b // ending with b

&> print and sent errors

